
Consumer ABS (Name WIP)

James T. Bishop
Department of Computer Science
Boston University
Boston, MA 02215
jtbishop@bu.edu

Abstract

This paper explores statistical relationships in Consumer Asset Backed Securities including the pricing of individual securities and probability of prepayment or default. As with Mortgage-Backed Securities, the largest risk an investor in Consumer ABS faces is prepayment risk. Identifying relevant factors to accurately quantify this risk is paramount to extracting value from the Consumer ABS market.

Introduction

Asset Backed Securities (ABS) are a class of financial structured product that pool together cash flow generating assets and payout that cash to the securities' holders. Consumer ABS are a category of ABS where the underlying asset pools come from credit card loans, student loans, auto loans, or specific merchant loans like Affirm or Afterpay.

What makes Consumer ABS specifically interesting is how underdeveloped the market is, even relative to the rest of the ABS market or rest of the structured product market. Of the \$9.2 trillion structured products market, non-mortgage ABS comprised \$1.3 trillion, with consumer ABS making up about \$500 billion of this .

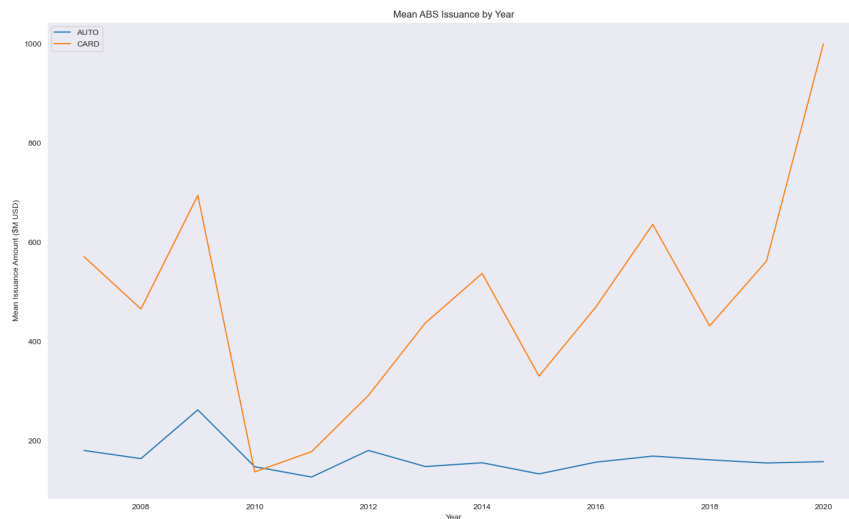


Figure 1: Mean Auto and Card ABS Issuance

Pricing these securities is difficult for a few reasons. First, while many of these securities trade frequently, there is less transparency than in equity and other fixed income markets. Many consumer ABS deals don't have public information available, and even those that do aren't necessarily part of FINRA's Trade Reporting and Compliance Engine (TRACE), which functions to report trade data similar to what's available to the general public on stock trades. However, this data isn't typically available to retail investors because of a paywall.

Second, most retail investors can't trade consumer ABS through their broker. This means fewer market participants and, as a result, more mispricings.

Finally, ABS are complex products. Modeling their cash flows, prepayment and credit risk, and volatility is difficult, even if one assumes a perfectly efficient market. Due to this, there's a lower level of consistency in price determination between market makers (e.g. banks) which also makes price determination difficult.

The goal of this project is to act as a primer on using modern data science methods to explore the ABS market and make ballpark price estimations.

Data

Data was gathered through Boston University's Bloomberg Terminal subscription as well as the Federal Reserve of St. Louis (FRED) online database.

From Bloomberg, the SRCH <GO> function was used to screen for public ABS deals (i.e. ABS issuances involving publicly traded companies) transacted in US dollars and with collateral in the US. Later, the SRCH <GO> function was also used to find the TRACE eligible subset of those deals. Additionally, individual bond prices for 20 bonds was gathered along with collateral for those bonds. The pricing data was gathered on a daily frequency while the collateral data was gathered at monthly frequency.

From FRED, historic yields for the 1, 2, 3, 5, 7, 10, 20, and 30 year treasuries was pulled on a monthly basis.

For data maps, see appendix.

Implementation

To start, I explored the data gathered using SECF <GO> and examined year over year mean issuance by ABS category, distribution of issuances by coupon type within each group and number and average dollar amount of each issuance by coupon type. Next, I removed non-TRACE eligible bonds from my dataset and looked at a correlation matrix of the remaining securities' features. I then used K-Means clustering with 10 clusters to find 10 general classes of TRACE eligible bonds, based on available features. From here, I took note of the 10 cluster centers and randomly selected one other security from a cluster. I used these 20 TRACE eligible bonds later for pricing and collateral. I made this decision since there was no feasible way for me to fetch data on a massive bond universe given the data restrictions imposed by Boston University's Bloomberg subscription.

Following exploration, I made loader objects to load in my various data sources. These are all in the `preprocessing.py` file. I then began to look at ways to predict prepayments on the bonds in the TRACE Eligible dataset. I included all feature columns except "Is Mortgage Paid Off", "Next Call Date", "WAC", "Current WAL", and "Amt Out" because "Is Mortgage Paid Off" is our target variable, and the remaining variables' conditional distributions would reveal very clearly whether or not a security had been prepaid. I then took use One Hot Encoding on the categorical features, made the date features Gregorian ordinal values, and used a label encoder on the remaining columns.

I trained kNN, SVM, Gaussian Process, Decision Tree, Random Forest Classifier, Multilayer Perceptron, AdaBoost, Naive Bayes, and Quadratic Discriminant classifiers on 80% of the data and tested them on 20%. Results will be discussed in .

Following the classification task, I attempted to use Generalized Least Squares to predict monthly returns on the 20 securities mentioned earlier. I merged the price and collateral DataFrames with

historic treasury yields. Since this was timeseries data, I then lagged the exogenous variables by 1 timestep.

Results

In classification, I considered RMSE, accuracy, recall, and F1 Score and found that a Decision Tree with a max depth of 14 performed best:

```
Test: RMSE of Pipeline(steps=[('decisiontreeclassifier',
                                DecisionTreeClassifier(max_depth=14))]) is 0.139
precision    recall  f1-score   support

      0       0.99      0.99      0.99        178
      1       0.93      0.93      0.93         29

 accuracy          0.98          207
 macro avg          0.96          207
weighted avg          0.98          207
```

In the return prediction task, results for Auto and Consumer loans looked like this:

Results for CARMX 2018-1 A4

```
GLS Regression Results
=====
Dep. Variable:          Price      R-squared (uncentered):      0.829
Model:                  GLS       Adj. R-squared (uncentered):    0.712
Method:                 Least Squares  F-statistic:              7.076
Date:                  Fri, 27 Nov 2020  Prob (F-statistic):      8.33e-05
Time:                  09:44:51      Log-Likelihood:           179.65
No. Observations:      32          AIC:                        -333.3
Df Residuals:          19          BIC:                        -314.3
Df Model:              13
Covariance Type:       nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
WAC          -0.0247      0.016      -1.564      0.134      -0.058      0.008
WAM           0.0027      0.001       2.775      0.012       0.001      0.005
WALA          0.0029      0.002       1.331      0.199      -0.002      0.008
Balance       1.9e-12    4.49e-11      0.042      0.967     -9.2e-11     9.58e-11
Principal    -8.48e-11    1.64e-10     -0.519      0.610     -4.27e-10     2.57e-10
DGS1          0.0020      0.011       0.178      0.861     -0.022      0.026
DGS2         -0.0182      0.055     -0.333      0.742     -0.133      0.096
DGS3          0.0095      0.069       0.137      0.892     -0.135      0.154
DGS5          0.0032      0.041       0.077      0.940     -0.083      0.089
DGS7         -0.0159      0.033     -0.483      0.635     -0.085      0.053
DGS10         0.0405      0.024       1.690      0.107     -0.010      0.091
DGS20        -0.0339      0.030     -1.139      0.269     -0.096      0.028
DGS30         0.0139      0.022       0.624      0.540     -0.033      0.060
=====
Omnibus:              1.544    Durbin-Watson:              2.084
Prob(Omnibus):        0.462    Jarque-Bera (JB):           1.018
Skew:                 -0.437    Prob(JB):                   0.601
Kurtosis:             2.987    Cond. No.                   4.15e+11
=====
```

Results for Card loans looked like:

Results for AMXCA 2018-3 A

GLS Regression Results

```

=====
Dep. Variable:          Price      R-squared (uncentered):          0.804
Model:                  GLS        Adj. R-squared (uncentered):          0.544
Method:                  Least Squares      F-statistic:                  3.084
Date:                   Fri, 27 Nov 2020    Prob (F-statistic):          0.0271
Time:                   09:44:51          Log-Likelihood:              181.82
No. Observations:       28              AIC:                        -331.6
Df Residuals:           12              BIC:                        -310.3
Df Model:               16
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
ExcessSpread3M	-0.0004	0.001	-0.566	0.582	-0.002	0.001
ExcessSpread1M	3.192e-05	0.001	0.053	0.959	-0.001	0.001
Mth Pay Rate	2.215e-05	0.000	0.141	0.891	-0.000	0.000
Port Yld	0.0002	0.000	0.473	0.645	-0.001	0.001
Chg Offs	-7.907e-05	0.001	-0.061	0.952	-0.003	0.003
30	-0.0128	0.011	-1.206	0.251	-0.036	0.010
60	0.0036	0.016	0.229	0.822	-0.031	0.038
del90Plus	0.0063	0.008	0.825	0.426	-0.010	0.023
DGS1	0.0022	0.009	0.244	0.811	-0.018	0.022
DGS2	-0.0262	0.037	-0.701	0.497	-0.108	0.055
DGS3	0.0288	0.050	0.576	0.575	-0.080	0.138
DGS5	0.0130	0.027	0.476	0.643	-0.047	0.073
DGS7	-0.0217	0.031	-0.705	0.494	-0.089	0.045
DGS10	0.0002	0.021	0.009	0.993	-0.045	0.046
DGS20	-0.0015	0.019	-0.075	0.941	-0.044	0.041
DGS30	0.0054	0.021	0.262	0.798	-0.040	0.051
=====						
Omnibus:	14.666	Durbin-Watson:	2.219			
Prob(Omnibus):	0.001	Jarque-Bera (JB):	26.520			
Skew:	0.924	Prob(JB):	1.74e-06			
Kurtosis:	7.395	Cond. No.	4.49e+04			
=====						

Both sets of data showed statistical significance at at least the 5% level, based on simultaneous hypothesis tests for the explanatory variables in the model.

Challenges

The biggest challenge in this project was data availability. Despite the cleanliness of data provided by Bloomberg, the amount of data with useful features was limited. It was impossible to pull a large universe of securities all at once, even in the screening phase which meant that there were approximately 1,000 securities used in the classification task. Similarly, it was impossible to fetch a high volume of individual bond prices or collateral at once, and scripting was not an option which meant a small universe of securities for tasks using that data.

Another evident challenge is the complexity of ABS. For instance, the LIBOR Market Model (LMM) is a popular tool for pricing structured products, among other things. However, calibrating the model is a complex optimization process as is applying the model to data, particularly when the data available is low frequency and there isn't a lot available.

Conclusion & Takeaways

The largest takeaway is that a Decision Tree Classifier is very well suited to the task of predicting whether a security has defaulted or not, given this dataset. Obviously, one key issue with this is that we haven't predicted when or how likely that default is, which means that it's difficult to quantify the risk associated with it. This is discussed in greater detail in .

This paper also provides a starting point for time series analysis of ABS prices. By looking at a small number of factors, relatively accurate predictions were able to be made. To really see the performance of the GLS model, however, backtesting would be required, which wasn't possible with the data available. This is also a source of future work.

Future Work

I plan to gather incrementally more data and build a larger set of securities with collateral and pricing. In doing so, I'll be able to more produce results with a higher degree of certainty.

I also intend to examine price at issue data and gather text files of filing documents and news on companies issuing these debt securities to see whether or not sentiment analysis has a place in accurately pricing initial offerings.

Further, I am working on implementing LMM. The difficulty lies in the frequency of data I have available because typically the model is used on extremely high frequency data, e.g. 1 minute intervals rather than days or months.

As mentioned earlier, the Decision Tree Classifier doesn't answer "How risky is this security?" or "When will we likeliest see prepayments occur?". To do so requires a more complex model, and is something I plan to explore in the future, hopefully using properties of the LMM model.

Appendix

Table 1: ABS Trace Issuance Datasets

Label	Type	Description
Amt Out	float64	Amount outstanding (USD)
BBG Composite	object	Bloomberg Composite credit rating
CUSIP	object	
Cpn	float64	Coupon (%)
Current WAL	float64	Current Weighted Average Life
Day Count	object	Day count convention (30/360, 30/365, or Actual)
Delinquency Rate 60+ Days	float64	
Delinquency Rate 90+ Days	float64	
Issue Date	datetime64[ns]	
Issuer Name	object	
Maturity	datetime64[ns]	
Mid Price	float64	Mid price between bid and ask price
Mortgage Original Amount	int64	
Next Call Date	datetime64[ns]	Only applicable to callable bonds
Next Coupon Date	datetime64[ns]	
Original Maximum Loan Size	float64	Size of largest loan in the pool at issue
Price at Issue	float64	
Security Name	object	
Ticker	object	
Category	object	CARD, AUTO, or CONSUMER
Original WAL	float64	Original Weighted Average Life
isCallable	float64	

Table 2: TRACE Eligible Loans Dataset

Label	Type	Description
CUSIP	object	
Security Name	object	
Mortgage Original Amount	float64	
Cpn	float64	Coupon (%)
Current WAL	float64	
Amt Out	float64	Amount outstanding (USD)
BBG Composite	object	Bloomberg Composite credit rating
Day Count	object	Day count convention (30/360, 30/365, or Actual)
Delinquency Rate 60+ Days	float64	
Delinquency Rate 90+ Days	float64	
Is Mortgage Paid Off	int64	
Issue Date	datetime64[ns]	
Maturity	datetime64[ns]	
Next Call Date	datetime64[ns]	
Ticker	object	
Price at Issue	float64	
Benchmark Spread at Issue	float64	Spread to the benchmark rate at issue
PSA Since Issuance	float64	Public Securities Association prepayment calculation
WAC	float64	Weighted Average Coupon
Category	object	CARD, AUTO, or CONSUMER
isCallable	float64	
Original WAL	float64	Original Weighted Average Life

Table 3: Collateral for Credit Card ABS

Label	Type	Description
ExcessSpread3M	float64	
ExcessSpread1M	float64	
Mth Pay Rate	float64	
Port Yld	float64	
Chg Offs	float64	
30	float64	30 day delinquency rate
60	float64	60 day delinquency rate
del90Plus	float64	90 day delinquency rate

References

Table 4: Collateral for Auto and Consumer ABS

Label	Type	Description
WAC	float64	Weighted Average Coupon
WAM	int64	Weighted Average Maturity
WALA	int64	Weighted Average Loan Age
Balance	int64	
Principal	float64	

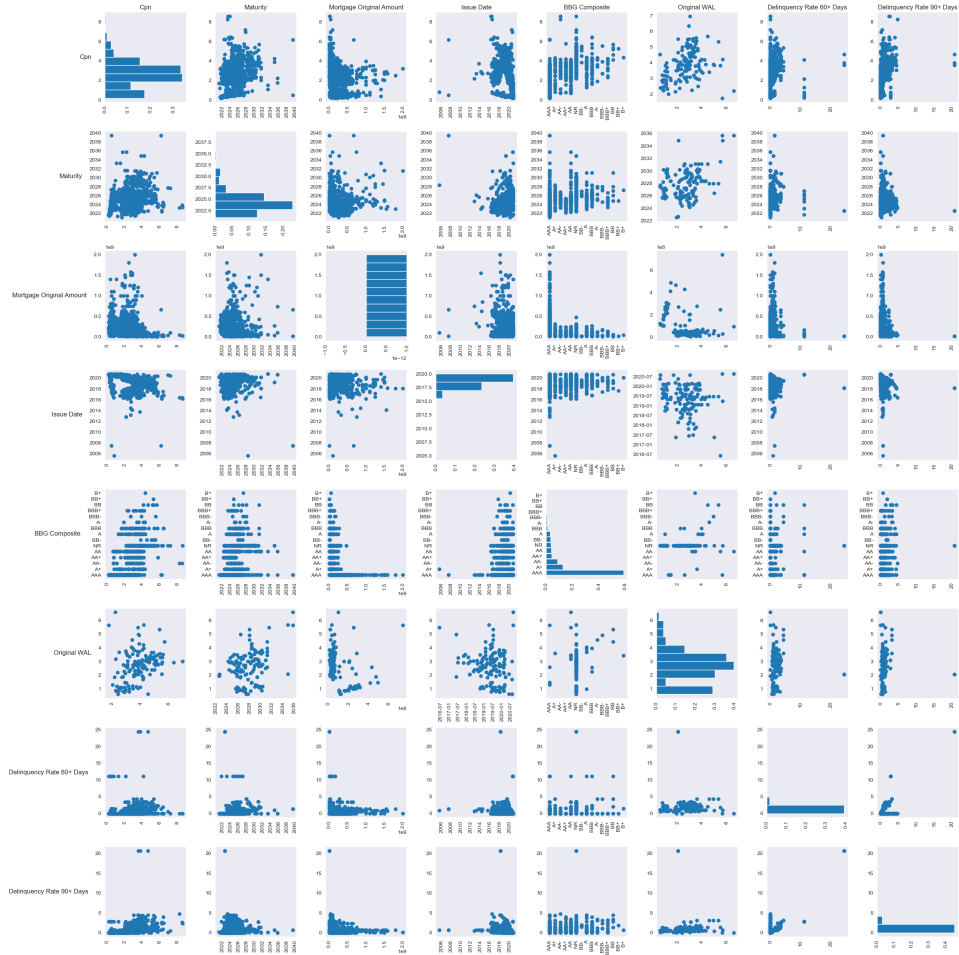


Figure 2: Correlation Matrix of TRACE Eligible Bonds